

Attention kernel latency

One H100 SXM5. Same machine, same harness, same shapes, and every point has a passing parity run behind it. The hollow marker is biasless attention: the floor this feature can approach but never beat. Speedups are quoted against `score_mod`, the path vLLM actually ships, or against the fastest day-0 path measured where `score_mod` was not run.

