


A Attention kernel latency, one H100 SXM5

Same machine, same harness, same shapes. Every timing has a passing parity run behind it. Grey is biasless attention: the floor this feature can approach but never beat.



B Status by unit and architecture

Green means it ran on that silicon and passed. Nothing is marked done on hardware it has never run on.

	sm_120 RTX 5090	sm_90 H100	sm_80 A100	sm_100 B200
U2 relative attention	validated	validated	validated	no hardware
U2 full-model token gate	n/a	32 of 32	n/a	no hardware
U3 FP8 paged KV, writes	validated	validated	validated	no hardware
U3 FP8 paged KV, reads	pending	pending	pending	pending
U2 split-KV decode	pending	pending	pending	pending
End-to-end serving curves	n/a	run lost	n/a	pending
U1 MoE grouped GEMM	pending	pending	pending	pending



validated on that silicon



ran, results lost



not done, or not applicable