


A Attention kernel latency

One H100 SXM5. Same machine, same harness, same shapes; every point has a passing parity run behind it. The hollow marker is biasless attention, the floor this feature can approach but never beat. Speedups are quoted against score_mod, the path vLLM actually ships, or against the fastest day-0 path measured where score_mod was not run.



B Validation status

Filled means it ran on that silicon and passed its gate. Nothing is marked done on hardware it has never run on.

	sm_120 RTX 5090	sm_90 H100	sm_80 A100	sm_100 B200
U2 relative attention, global and SWA	●	●	●	○
U2 full-model token gate, 32 prompts	—	●	—	○
U3 FP8 paged KV, write path	●	●	●	○
U3 FP8 paged KV, read path	○	○	○	○
U2 split-KV decode	○	○	○	○
End-to-end serving curves	—	○	—	○
U1 MoE grouped GEMM	○	○	○	○

● validated on that silicon ○ ran, results lost ○ not done yet — not applicable